


TABLA: A Framework for Accelerating Statistical Machine Learning

Presenters:

MeiXing Dong, Lajanugen Logeswaran

Intro

- Machine learning algorithms widely used, computationally intensive
- FPGAs get performance gains w/ flexibility
- Development for FPGAs expensive and long
- Automatically generate accelerators (TABLA)





ISTOCK/ANNA LURYE

CAT

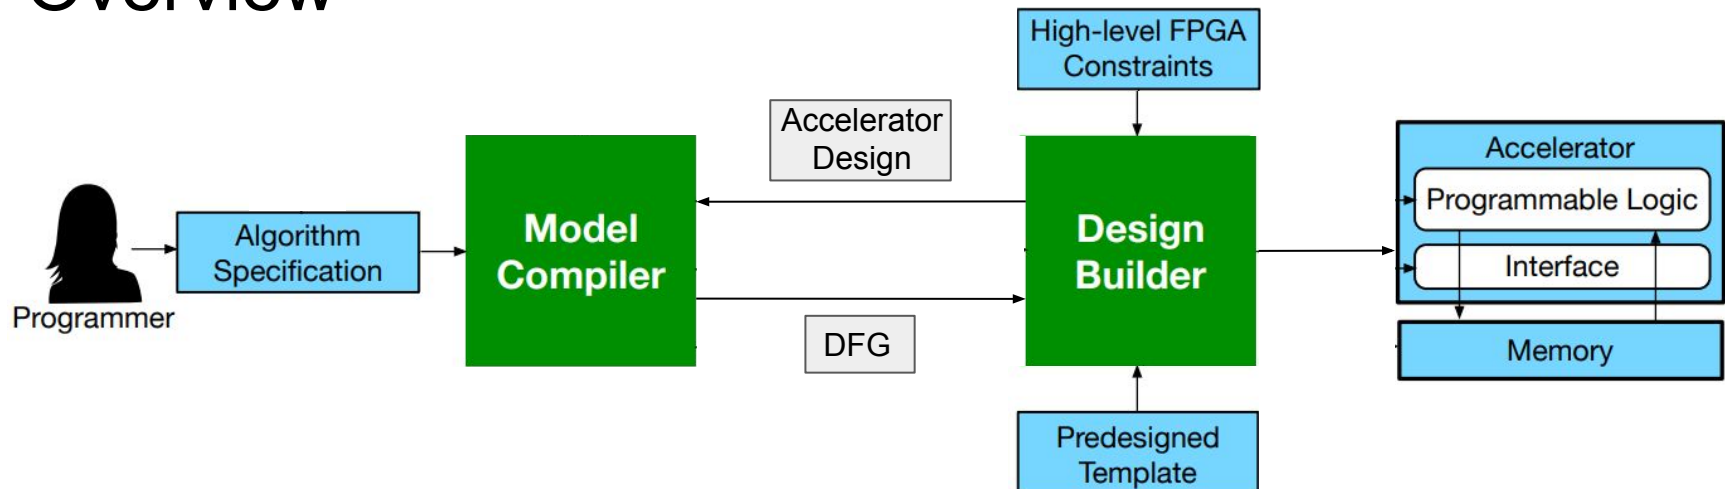
Stochastic Gradient Descent

- Machine learning uses objective (cost) functions
- Ex. linear regression
 - objective: $\sum_i 1/2(w^T x_i - y_i)^2 + \lambda ||w||$
 - gradient: $\sum_i (w^T x_i - y_i) x_i + \lambda ||w||$
- Want to find lowest value possible w/ gradient descent
- Can approximate batch update



Src: <https://alykhantejani.github.io/a-brief-introduction-to-gradient-descent/>

Overview





```

model_input  x[m]; //model input features
model_output y'[n]; //model outputs
model        w[n][m]; //model parameters
gradient      g[n][m]; //gradient

iterator i[0:m]; //iterator for group operations
iterator j[0:n]; //iterator for group operations

//m parallel multiplications followed by
//an addition tree; repeat n times in parallel
s[j] = sum[i](x[i] * w[j][i]);

y[j] = sigmoid(s[j]); //n parallel sigmoid operations
e[j] = y[j] - y'[j]; //n parallel subtractions
g[j][i] = x[i] * e[j]; //n*m parallel multiplications
rg[j][i] = λ * w[i][j]; //n*m parallel multiplications
g[j][i] = g[j][i] + rg[j][i]; //n*m parallel additions
    
```



Programming Interface

- Language

- Close to mathematical expressions
- Language constructs commonly used in ML algorithms

Type	Classification	Language Keywords
Data	Model inputs	model_input
	Model outputs	model_output
	Model Parameters	model
	Gradient of objective function	gradient
	Iterator variable	iterator
Operation	Basic	+, -, <, >, *
	Group	pi, sum, norm
	Non Linear	gaussian, sigmoid, sigmoid_symmteric, log

```
model_input  x[m]; //model input features
model_output y'[n]; //model outputs
model        w[n][m]; //model parameters
gradient     g[n][m]; //gradient

iterator i[0:m]; //iterator for group operations
iterator j[0:n]; //iterator for group operations


//m parallel multiplications followed by
//an addition tree; repeat n times in parallel
s[j] = sum[i](x[i] * w[j][i]);

y[j] = sigmoid(s[j]); //n parallel sigmoid operations
e[j] = y[j] - y'[j]; //n parallel subtractions
g[j][i] = x[i] * e[j]; //n*m parallel multiplications
rg[j][i] =  $\lambda$  * w[i][j]; //n*m parallel multiplications
g[j][i] = g[j][i] + rg[j][i]; //n*m parallel additions
```


- Why not MATLAB/R ?

- Identifying parallelizable code
- Conversion to hardware design

Model Compiler



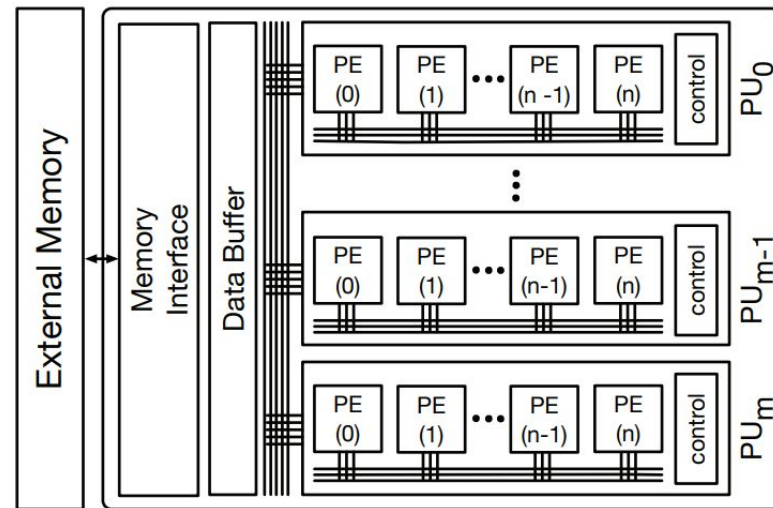
- Model parameters and gradient are both arrays of values
- Gradient function specified using math
- Ex.
 - $g[j][i] = u * g[j][i]$
 - $g[j][i] = w[j][i] - g[j][i]$



- Minimum-Latency Resource Constrained Scheduling
- Priority placed on highest distance from sink
- Predecessors scheduled
- Resources available

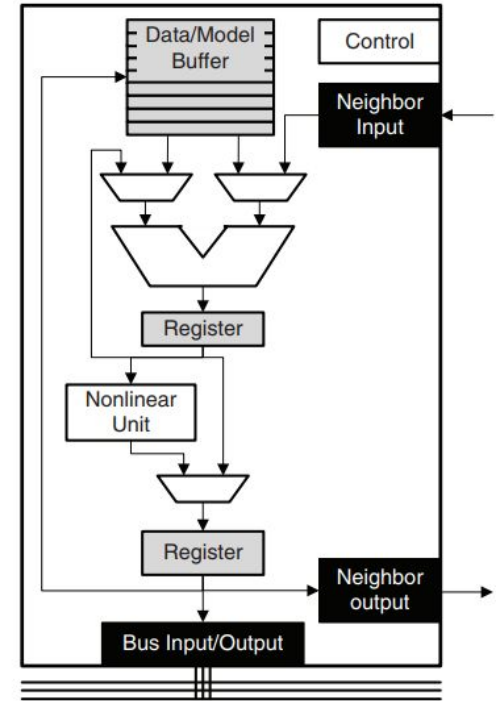
Accelerator Design: Design builder

- Generates Verilog of accelerator from
 - DFG, algorithm schedule, FPGA spec
- Clustered hierarchical architecture
- Determines
 - Number of PEs
 - Number of PEs per PU
- Generate
 - Control units and buses
 - Memory interface unit and access schedule



Accelerator Design: Processing engine

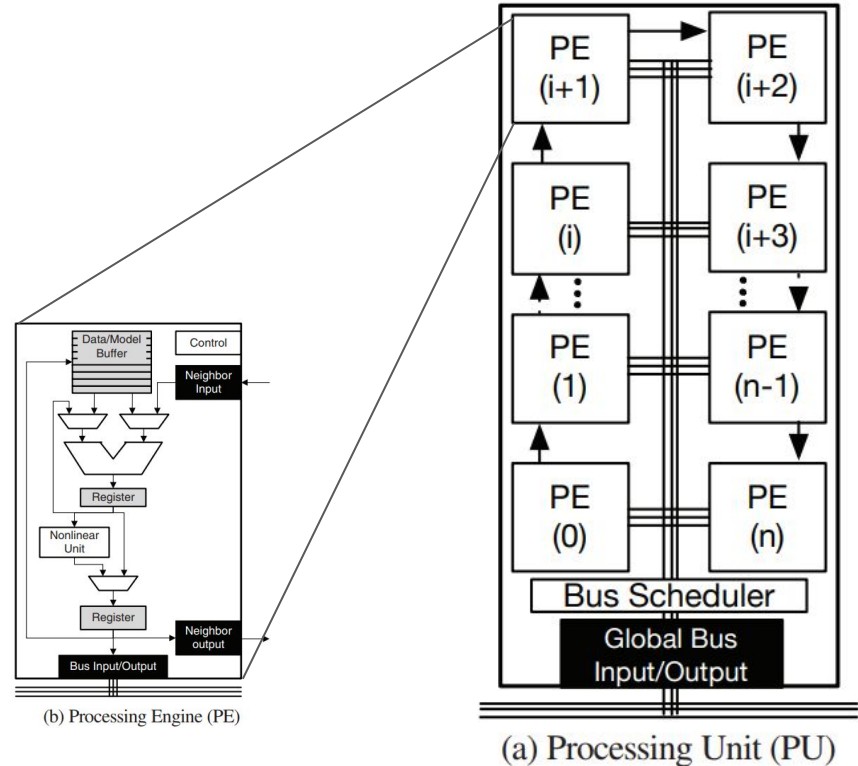
- Basic block
- Fixed components
 - ALU
 - Data/Model buffer
 - Registers
 - Busing logic
- Customizable components
 - Control unit
 - Nonlinear unit
 - Neighbor input/output communication



(b) Processing Engine (PE)

Accelerator Design: Processing unit

- Group of PEs
 - Modular design
 - Data traffic locality within PU
- Scale up as necessary
- Static communication schedule
 - Global bus
 - Memory access




Evaluation

Setup


- Implement TABLA using off-the-shelf FPGA platform (Xilinx Zynq ZC702)
- Compare with CPUs and GPUs
- 5 popular ML algorithms
 - Logistic Regression
 - Support Vector Machines
 - Recommender Systems
 - Backpropagation
 - Linear Regression
- Measurements
 - Execution time
 - Power

Performance Comparison



(b) Speedup of GPUs and TABLA design in comparison ARM A15.

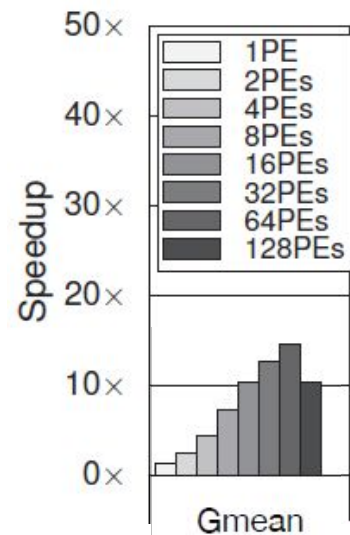
Power Usage



(b) Performance-per-Watt comparison between Tegra, GTX 650 , Tesla and TABLA

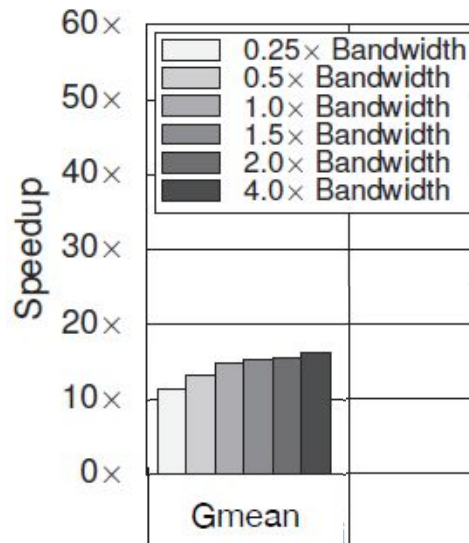
Design Space Exploration

- Number of PEs vs PUs
 - Configuration that provides highest frequency
 - 8 PEs per PU
- Number of PEs
 - Initially linear increase
 - Poor performance after a certain point
- Too many PEs
 - Wider global bus - Reduced frequency



Design Space Exploration

- Bandwidth sensitivity
 - Increase bandwidth between external memory and accelerator
 - Limited improvement
 - Computation dominates execution time
 - Frequently accessed data are kept in PE's local buffers



Conclusion

- Machine learning algorithms popular but compute-intensive
- FPGAs are appealing for accelerating performance
- FPGA design long and expensive
- Automatically generate accelerators for learning algorithms using template-based framework (TABLA)

Discussion Points

- Is this more useful than accelerators specialized for gradient descent?
- Is this solution practical? (Cost, Scalability, Performance)
- Is this idea generalizable to problems other than gradient descent?